

Chapter 6: Database Design using ERD ¹

Abu Raihan Mostofa Kamal

Professor, CSE Department
Islamic University of Technology (IUT)

September 12, 2023

¹Slides are based on Textbook, its companion slide and other sources

Chapter Outline

Overview of the Design Process

ER Model

Mapping Cardinality

Generalization/Specialization

Design Phases

1. **The initial phase:** Characterize fully the data needs of the prospective database users. The database designer needs to interact extensively with **domain experts and users** to carry out this task. The outcome of this phase is a **specification of user requirements**.
2. **Conceptual-design phase:** The **entity-relationship model** is typically used to represent the conceptual design. This is an abstract-design of the system. This is **independent** of any specific database implementation.
3. **Implementation phase:** It has two levels:
 - (a) **Logical Design:** It maps the high-level conceptual schema onto the **implementation data model of the database system**.
 - (b) **Physical Design:** It takes care of the **actual storage features** such as default space, indexing and so on.



Design Phases

1. **The initial phase:** Characterize fully the data needs of the prospective database users. The database designer needs to interact extensively with **domain experts and users** to carry out this task. The outcome of this phase is a **specification of user requirements**.
2. **Conceptual-design phase:** The **entity-relationship model** is typically used to represent the conceptual design. This is an abstract-design of the system. This is **independent** of any specific database implementation.
3. **Implementation phase:** It has two levels:
 - (a) **Logical Design:** It maps the high-level conceptual schema onto the implementation data model of the **database system**.
 - (b) **Physical Design:** It takes care of the **actual storage features** such as default space, indexing and so on.



Design Phases

1. **The initial phase:** Characterize fully the data needs of the prospective database users. The database designer needs to interact extensively with **domain experts and users** to carry out this task. The outcome of this phase is a **specification of user requirements**.
2. **Conceptual-design phase:** The **entity-relationship model** is typically used to represent the conceptual design. This is an abstract-design of the system. This is **independent** of any specific database implementation.
3. **Implementation phase:** It has two levels:
 - (a) **Logical Design:** It maps the high-level conceptual schema onto the **implementation** data model of the **database system**.
 - (b) **Physical Design:** It takes care of the **actual storage features** such as default space, indexing and so on.

Design Phases

1. **The initial phase:** Characterize fully the data needs of the prospective database users. The database designer needs to interact extensively with **domain experts and users** to carry out this task. The outcome of this phase is a **specification of user requirements**.
2. **Conceptual-design phase:** The **entity-relationship model** is typically used to represent the conceptual design. This is an abstract-design of the system. This is **independent** of any specific database implementation.
3. **Implementation phase:** It has two levels:
 - (a) **Logical Design:** It maps the high-level conceptual schema onto the **implementation** data model of the **database system**.
 - (b) **Physical Design:** It takes care of the **actual storage features** such as default space, indexing and so on.

Design Phases

1. **The initial phase:** Characterize fully the data needs of the prospective database users. The database designer needs to interact extensively with **domain experts and users** to carry out this task. The outcome of this phase is a **specification of user requirements**.
2. **Conceptual-design phase:** The **entity-relationship model** is typically used to represent the conceptual design. This is an abstract-design of the system. This is **independent** of any specific database implementation.
3. **Implementation phase:** It has two levels:
 - (a) **Logical Design:** It maps the high-level conceptual schema onto the **implementation** data model of the **database system**.
 - (b) **Physical Design:** It takes care of the **actual storage features** such as default space, indexing and so on.

Bad and Good Design

In designing a database schema, we must ensure that we avoid **two major problems**:

1. **Redundancy**: A bad design may repeat information. The **biggest problem** with such redundant representation of information is that the copies of a piece of information can become **inconsistent** if the information is updated.
2. **Incompleteness**: A bad design may make certain aspects of the enterprise difficult or impossible to model. It introduces **bad business logic**.



Design Approaches

- **Entity Relationship Model** (will be covered in this chapter).
 - ✓ Collection of **entities and relationships**
 - ✓ Entity is an object (with a number of attributes), while relationship defines how entities are related.
 - ✓ Represented **diagrammatically** by an entity-relationship diagram
- **Normalization Theory:** Formal method to test if the design is good or bad. (will be covered in Chapter 7)

Overview of ER Model

Entity Sets

- **Entity:** An entity is an object that **exists** and is **distinguishable** from other objects.
Example: specific person, company, event, plant
- An entity set is a set of **entities of the same type that share the same properties**.
Example: set of all persons, companies, trees, courses
- An entity is represented by a **set of attributes** i.e., descriptive properties possessed by all members of an entity set.

Example:

```
instructor = (ID, name, salary )  
course= (course_id, title, credits)
```

- A subset of the attributes form a **primary key** of the entity set; i.e., uniquely identifying each member of the set.

Entity sets in ER Diagram

Entity sets can be represented graphically as follows:

- Rectangles represent entity sets
- Attributes listed inside entity rectangle
- Underline indicates primary key attributes




Figure: E-R diagram for entity sets instructor and student

Relationship Sets

- A relationship is an association among several entities

44553 (Peltier) (student entity)

advisor (relationship set)

22222 (Einstein) (instructor entity)

- A relationship set is a mathematical relation among $n \geq 2$ entities, each taken from entity sets

$$\{(e_1, e_2, \dots, e_n) | e_1 \subseteq E_1, e_2 \subseteq E_2, \dots, e_n \subseteq E_n\}$$

where (e_1, e_2, \dots, e_n) is a relationship.

Example: $(44553, 22222) \subseteq \text{advisor}$

Relationship Sets (Cont.)

- **Example:** we define the relationship set advisor to denote the associations between students and the instructors who act as their advisors.
- **Pictorially,** we draw a line between related entities as follows:




Figure: Relationship set: Example

Represent Relationship Sets

- **Diamonds** represent relationship sets.




Figure: Diamond is the Relationship symbol

Relationship Sets: Additional Attribute

- An **attribute** can also be **associated** with a relationship set.
- For instance, the advisor relationship set between entity sets instructor and student may have the **attribute** **date** which tracks when the student started being associated with the advisor

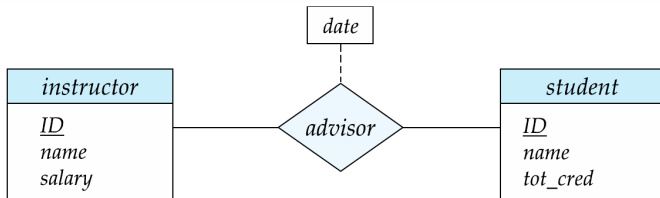


Figure: Additional Attribute in the Relationship

Relationship Sets: Roles

- Entity sets of a relationship need not be distinct, each relation may be used multiple times with different roles.
- For instance, `course_id` and `prereq_id` are called **roles**
- In fact, here same attribute (i.e. `course_id`) is used for 2 purposes.

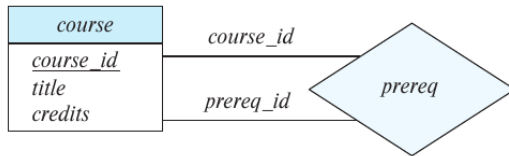


Figure: Roles in the Relationship

Relationship Sets: Roles (Cont.)

- In this type of relationship set, sometimes called a **recursive** relationship set.
- The records of one entity refers to itself. For instance, course entity here will have at least two attributes: (i) **CourseID** and (ii) **PreReqID** .
- In actual implementation, such **Roles** are achieved by **self-referencing**.

Degree of a Relationship Set

- The number of entity sets that participate in a relationship set is the **degree** of the relationship set.
- **Most** of the relationship sets in a database system are **binary**.
- Occasionally, however, relationship sets involve more than two entity sets.
- **Example:** Consider the entity projects that represents all the research projects of the university. Each project can have multiple associated students and multiple associated instructors. It becomes a **Ternary Relationship**.




Figure: 3 Entities in Relationship

Complex Attributes

In the E-R model, attributes can be grouped into as follows:

- Simple and composite attributes
- Single-valued and multivalued attributes.
- Derived attributes.



Simple and Composite Attributes

- Simple attribute has **no subparts**. While Composite attribute **has subparts** (i.e. Name= first name, last name).




Figure: Composite Attribute

- Composite attribute may appear as a **hierarchy**. (each level is called component attribute)
- Composite attribute is **preferable** if different components are useful at different occasions. **For instance**, if we need some report of employees according to last name for HR section, while another list with first name is needed for accounts section.

Simple and Composite Attributes: Implementation

- In Oracle Database, it is implemented as **TYPE** as **OBJECT** as a user-defined datatype containing sub-parts.
- Oracle Database requires you to use a **table alias** to qualify any dot-notational reference to subprograms or attributes of objects.
- Use of a table **alias is optional when referencing top-level attributes** of an object table directly, without using the dot notation.



Single-valued and Multi-valued attributes

- The attributes in our examples all have a single value for a particular entity. For instance, the student ID attribute for a specific student entity refers to only one student ID. Such attributes are said to be single valued.
- An attribute may have zero, one or more values against one record, then it is called multi-valued attribute. Example: one citizen may have more than one phone number.



Derived attributes

- The value for this type of attribute can be **derived from the values of other** related attributes or entities.
- **For instance**, age can be derived from date of birth.
- The value of a derived attribute is **not stored but is computed when required**.

3 categories presented are based on the concept of **Domain of attribute**, which defines the permissible values and number of values for that attribute.

Complex Attributes in ER Diagram

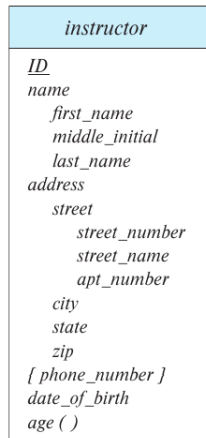


Figure: E-R diagram with composite, multivalued, and derived attributes

Composite Attributes: Implementation (Cont.)

1. Step 1: Create the type first (using Object concept)

```
create type nametype as object (  
  fname varchar2(10),  
  middlename varchar2(10),  
  lastname varchar2(10)  
);
```

2. Step 2: Create table with user-defined type as an attribute

```
create table emp  
  (id number primary key,  
   name nametype,  
   address varchar2(20),  
   constraint cnamenn check (name.fname IS NOT NULL)  
  );
```

Composite Attributes: Implementation (Cont.2)

1. Step 3: Data insert and select

```
insert into emp values (101,  
nametype('Abdur', 'Rahim', 'Mia'), 'Uttara');
```

```
% Select can be done this way:
```

```
select id, e.name.fname, address from emp e;
```

```
select id, name, address from emp;
```



Multi-valued Attributes: Implementation

1. Create attribute using **varray**, then create table and finally insert data

```
CREATE OR REPLACE TYPE vmobiles AS VARRAY(10) OF VARCHAR2(20);

CREATE TABLE students_multiple (id number, name varchar2(20), phone vmobiles);

insert into students_multiple values(1, a ,vmobiles( 0001 ));

insert into students_multiple values(2, b ,vmobiles( 0002 , 0003 , 0004 ));

insert into students_multiple values(3, c ,vmobiles(NULL));
```



Multi-valued Attributes: Implementation (Select)

1. You will have to use the TABLE operator in the case of a VARRAY
-

```
--this is how you need to select
```

```
--You will have to use the TABLE operator in the case of a VARRAY
```

```
select id,name, e.*  
FROM STUDENTS_MULTIPLE s, TABLE(s.phone) e;
```

```
select * from  
students_multiple;
```



Mapping Cardinality Constraints

- Express the number of entities to which another entity can be associated via a relationship set.
- Most useful in describing binary relationship sets.
- For a binary relationship set the mapping cardinality must be one of the following types:
 - (a) One to one
 - (b) One to many
 - (c) Many to one
 - (d) Many to many

Mapping Cardinality: One to One

- **One to one:** An entity in A is associated with at most one entity in B, and an entity in B is associated with at most one entity in A.

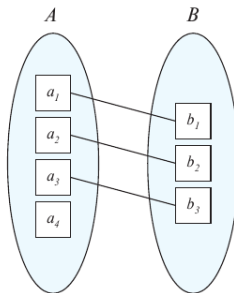


Figure: One to One Mapping

Mapping Cardinality: One to Many

- **One to Many:** An entity in A is associated with any number (zero or more) of entities in B. An entity in B, however, can be associated with at most one entity in A.




Figure: One to Many Mapping

Mapping Cardinality: Many to One

- **Many to One:** An entity in A is associated with at most one entity in B. An entity in B, however, can be associated with any number (zero or more) of entities in A.




Figure: Many to One Mapping

Mapping Cardinality: Many to Many

- **Many to Many:** An entity in A is associated with any number (zero or more) of entities in B, and an entity in B is associated with any number (zero or more) of entities in A.

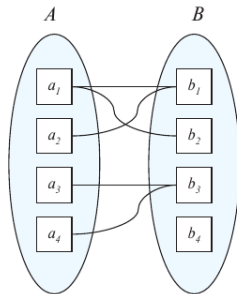


Figure: Many to Many Mapping

Cardinality Constraints in ER Diagram

- A directed line \longrightarrow is used to specify "one"
- An un-directed _____ line specify "many" which is either zero or more
- **Example:** One-to-one relationship between an instructor and a student.
A student is associated with at most one instructor via the relationship advisor

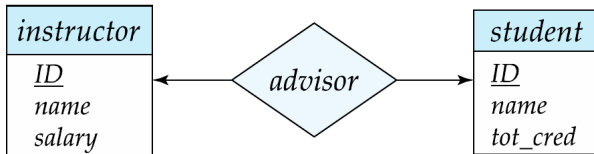


Figure: One to One Mapping

ER Diagram for One-to-Many

- One-to-Many relationship between an instructor and a student
 - ✓ an instructor is associated with several (including 0) students via advisor
 - ✓ a student is associated with at most one instructor via advisor

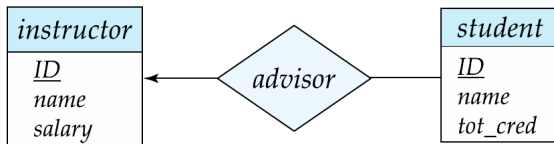


Figure: One to Many Mapping

ER Diagram for Many-to-One

- In a many-to-one relationship between an instructor and a student
 - ✓ an instructor is associated with at most one student via advisor
 - ✓ and a student is associated with several (including 0) instructors via advisor

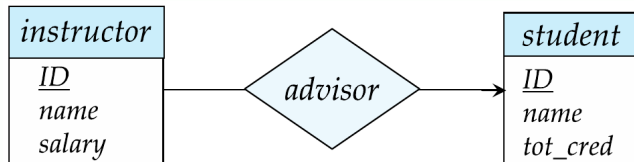


Figure: Many to One Mapping

ER Diagram for Many-to-Many

- Suppose we are allowing joint-supervision, students work in a group.
 - ✓ An instructor is associated with several (possibly 0) students via advisor
 - ✓ A student is associated with several (possibly 0) instructors via advisor

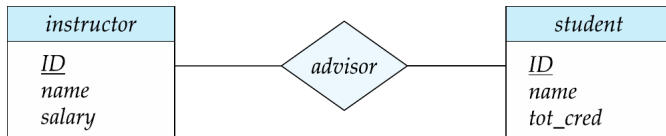


Figure: Many to Many Mapping

Total and Partial Participation

- **Total Participation¹ (indicated by double line):** every entity in the entity set participates in at least one relationship in the relationship set.

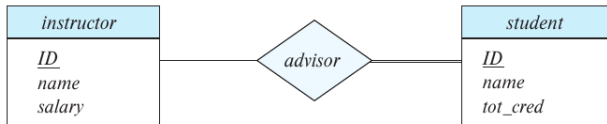


Figure: Total Participation

Example: participation of student in advisor relation is **total** which implies every student must have an associated instructor.

- Partial participation: some entities may not participate in any relationship in the relationship set.

Example: participation of instructor in advisor is partial which implies some instructors may not supervise any student

¹ For Total Participation NOT NULL constraint is used along with Foreign Key

Max..Min Cardinality: Alternative Representation

- E-R diagrams also provide a way to indicate more complex constraints on the **number of times each entity participates** in relationships in a relationship set.
- A line may have an associated minimum and maximum cardinality, shown in the form l..h, where l is the minimum and h the maximum cardinality.
- A minimum value of 1 indicates total participation of the entity set in the relationship set.
- 0..* means zero or more

Max..Min Cardinality: Precautions

- This min..max format could be **misleading** if not properly explained.

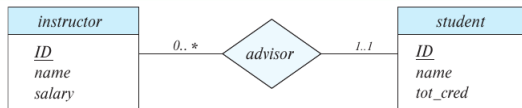


Figure: Min..Max format in cardinality

- The limit *0..** on the line between *advisor* and *instructor* indicates that an instructor can have zero or more students.
- The limit *1..1* on the line between *student* and *advisor* indicates that one student must have exactly one instructor as advisor.
- So, the relationship *advisor* is **one-to-many** from *instructor* to *student*. And participation of *student* in *advisor* is **total**.

Max..Min Cardinality: Precautions

- This min..max format could be **misleading** if not properly explained.

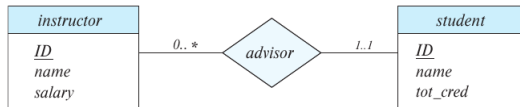


Figure: Min..Max format in cardinality

- The limit *0..** on the line between *advisor* and *instructor* indicates that **an instructor can have zero or more students**.
- The limit *1..1* on the line between *student* and *advisor* indicates that **one student must have exactly one instructor as advisor**.
- So, the relationship *advisor* is **one-to-many** from *instructor* to *student*. And **participation of student** in *advisor* is **total**.

Max..Min Cardinality: Precautions

- This min..max format could be **misleading** if not properly explained.




Figure: Min..Max format in cardinality

- The limit *0..** on the line between *advisor* and *instructor* indicates that **an instructor can have zero or more students**.
- The limit *1..1* on the line between *student* and *advisor* indicates that **one student must have exactly one instructor as advisor**.
- So, the relationship *advisor* is **one-to-many** from *instructor* to *student*. And **participation of student** in *advisor* is **total**.

Max..Min Cardinality: Precautions

- This min..max format could be **misleading** if not properly explained.




Figure: Min..Max format in cardinality

- The limit *0..** on the line between *advisor* and *instructor* indicates that **an instructor can have zero or more students**.
- The limit *1..1* on the line between *student* and *advisor* indicates that **one student must have exactly one instructor as advisor**.
- So, the relationship *advisor* is **one-to-many** from *instructor* to *student*. And **participation of student** in *advisor* is **total**.

Weak and Strong Entity Sets

- Strong Entity is independent of any other entity in the schema, in other words, it has sufficient attributes to form the primary key.
 - ✓ The strong entity is represented by a single rectangle.
 - ✓ The relationship between two strong entities is represented by a single diamond.
- A weak entity is an entity set that does not have sufficient attributes for Unique Identification of its records. (no primary key)



Weak Entity Sets

- It has **no primary key**. The existence of a weak entity set depends on the existence of a strong entity set
- It has **partial key**, which is combined with the primary key of the its strong entity set on the other-side of the relation to distinguish each record.
- Note: In most cases the **design is flawed**, a better design is possible where both will be strong entity sets.

Weak Entity Sets (Cont.)

- A **double rectangle** is used for representing a weak entity set
- The **double diamond** symbol is used for representing the relationship between a strong entity and a weak entity which is known as identifying relationship

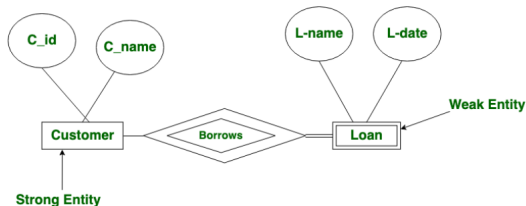


Figure: Weak Entity Set¹

¹image source:<https://www.geeksforgeeks.org/>

ER Diagram : Implementation

- In theory, there are 4 types of mappings, but we can eliminate one (since 1 to many and many to 1 becomes identical if we change the direction)
- First we will look at some **realistic examples** in each group before we implement them.
- Now the task goes in the following steps:



ER Diagram : Implementation

- In theory, there are 4 types of mappings, but we can eliminate one (since 1 to many and many to 1 becomes identical if we change the direction)
- First we will look at some **realistic examples** in each group before we implement them.
- Now the task goes in the following steps:



ER Diagram : Real-life Example

- One to One: Passport and Driving License
- One to Many: Department and students
- Many to Many: Students and Courses (taken)

ER Diagram: Implementation

- Suppose we have entity sets:
 - ✓ student (ID (pk), name, tot_cred, dept_name)
 - ✓ department (dept_name (pk), building, budget)
- We model the fact that each student has an associated department using a relationship set stud_dept
- The attribute dept_name in student below replicates information present in the relationship and is therefore **redundant** and it must be removed.
- BUT: when converting back to tables, in some cases the attribute gets reintroduced.

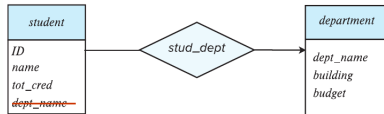


Figure: Redundant Attribute in ER-D

ER Diagram: Implementation (Cont.)

- **One to Many:** In many part use **foreign key** referencing the first entity.
- **One to One:** In any part (that depends on the context) use **foreign key with unique constraint**.
- **Many to Many:** Here a third entity (i.e. table) is needed termed as **Junction Table** which is **formed by two foreign keys of the entities**. Additional attributes may appear here.
- To ensure **Total Participation** in the first entity (for 1-1,1-m), additional mechanism is needed while, for **Total Participation in the second entity**, use **NOT NULL** constraint along with other constraint such as Foreign Key. (this principle is not applicable for m-m mapping)
- Remember: Total Participation may introduce additional business logic in many cases which may be impractical.

ER Diagram: Implementation- Practice

Consider the following scenario:

IUT has a number of departments. In each department, students are admitted. Each student must have ID, Name, Address. In each semester, a number of courses are offered. Each student can take a maximum of 48 courses.

- Your task is to draw the ER-D and implement it using DDLs.



Constraints on Generalization/Specialization in Database Design

Based on: Attribute of higher-level entity determines lower-level entity membership

- **Condition-defined:** Example: all customers over 65 years are members of senior-citizen entity set; senior-citizen 15% person. Since all the lower-level entities are evaluated on the basis of the same attribute (in this case, an age), this type of generalization is condition-defined.
- **User-defined:** Not constrained by a membership condition, rather, the designer simply assigns entities to a given entity set. For instance, let us assume that, after 3 months of employment, university employees are assigned to one of four work teams.



Constraints on Generalization/Specialization in Database Design

Based on: Attribute of higher-level entity determines lower-level entity membership

- **Condition-defined:** Example: all customers over 65 years are members of senior-citizen entity set; senior-citizen ISA person. Since all the lower-level entities are evaluated on the basis of the same attribute (in this case, on age), this type of generalization is said to be **attribute-defined**.
- **User-defined:** Not constrained by a membership condition, rather, the **database user** assigns entities to a given entity set. For instance, let us assume that, after 3 months of employment, university employees are assigned to one of four work teams.

Constraints on Generalization/Specialization in Database Design

Based on: Attribute of higher-level entity determines lower-level entity membership

- **Condition-defined:** Example: all customers over 65 years are members of senior-citizen entity set; senior-citizen ISA person. Since all the lower-level entities are evaluated on the basis of the same attribute (in this case, on age), this type of generalization is said to be **attribute-defined**.
- **User-defined:** Not constrained by a membership condition, rather, the **database user** assigns entities to a given entity set. For instance, let us assume that, after 3 months of employment, university employees are assigned to one of four work teams.



Constraints on Generalization/Specialization in Database Design

Based on: Attribute of higher-level entity determines lower-level entity membership

- **Condition-defined:** Example: all customers over 65 years are members of senior-citizen entity set; senior-citizen ISA person. Since all the lower-level entities are evaluated on the basis of the same attribute (in this case, on age), this type of generalization is said to be **attribute-defined**.
- **User-defined:** Not constrained by a membership condition, rather, the database user assigns entities to a given entity set. For instance, let us assume that, after 3 months of employment, university employees are assigned to one of four work teams.



Constraints on Generalization/Specialization in Database Design

Based on: Attribute of higher-level entity determines lower-level entity membership

- **Condition-defined:** Example: all customers over 65 years are members of senior-citizen entity set; senior-citizen ISA person. Since all the lower-level entities are evaluated on the basis of the same attribute (in this case, on age), this type of generalization is said to be **attribute-defined**.
- **User-defined:** Not constrained by a membership condition, rather, the **database user assigns** entities to a given entity set. **For instance**, let us assume that, after 3 months of employment, university employees are assigned to one of four work teams.



Constraints on Generalization/Specialization in Database Design

Based on: Attribute of higher-level entity determines lower-level entity membership

- **Condition-defined:** Example: all customers over 65 years are members of senior-citizen entity set; senior-citizen ISA person. Since all the lower-level entities are evaluated on the basis of the same attribute (in this case, on age), this type of generalization is said to be **attribute-defined**.
- **User-defined:** Not constrained by a membership condition, rather, the **database user assigns** entities to a given entity set. **For instance**, let us assume that, after 3 months of employment, university employees are assigned to one of four work teams.



Constraints on Generalization/Specialization (cont.)

Based on: The number of branching in its lower-level entity (i.e. leaf)

- **Disjoint.** A disjointness constraint requires that an entity belong to no more than one lower-level entity set. **For example**, Student entity can satisfy only one condition for the student type attribute; an entity can be either a graduate student or an undergraduate student, **but cannot be both**.
- **Overlapping.** The same entity **may belong to more than one lower-level entity set** within a single generalization. **For instance**, consider the employee work-team example, and assume that certain employees participate in more than one work team.

Constraints on Generalization/Specialization (cont.)

Based on: The number of branching in its lower-level entity (i.e. leaf)

- **Disjoint.** A disjointness constraint requires that an entity belong to no more than one lower-level entity set. **For example,** Student entity can satisfy only one condition for the student type attribute; an entity can be either a graduate student or an undergraduate student, **but cannot be both.**
- **Overlapping.** The same entity **may belong to more than one lower-level entity set** within a single generalization. **For instance,** consider the employee work-team example, and assume that certain employees participate in more than one work team.

Constraints on Generalization/Specialization (cont.)

Based on: Completeness

- **Total generalization or specialization.** Each higher-level entity (i.e. root) must belong to a lower-level entity set (i.e. leaf)
- **Partial generalization or specialization.** Some higher-level entities may not belong to any lower-level entity set

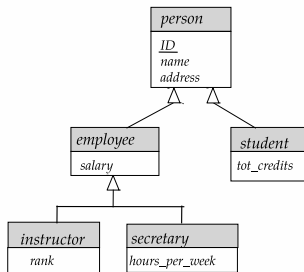
Constraints on Generalization/Specialization (cont.)

Based on: Completeness

- **Total generalization or specialization.** Each higher-level entity (i.e. root) must belong to a lower-level entity set (i.e. leaf)
- **Partial generalization or specialization.** Some higher-level entities may not belong to any lower-level entity set

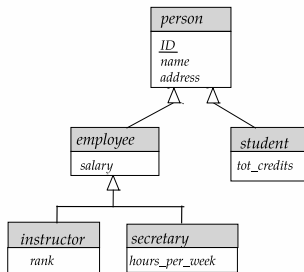
Representation of Generalization: 2 Ways

Consider the following E-R Diagram (ignore last 2 entities)



Representation of Generalization: 2 Ways

Consider the following E-R Diagram (ignore last 2 entities)



Representation of Generalization: Method 1

Method 1

- Create a schema for the higher-level entity set.
- For each lower-level entity set create it and link it with its upper-level schema.
- The primary-key attributes of the higher-level entity set become primary-key attributes of the higher-level entity set as well as all lower-level entity sets.
- We create foreign-key constraints on the lower-level entity sets.

Representation of Generalization: Method 1

Method 1

- Create a schema for the higher-level entity set.
- For each lower-level entity set create it and link it with its upper-level schema.
- The primary-key attributes of the higher-level entity set become primary-key attributes of the higher-level entity set as well as all lower-level entity sets.
- We create foreign-key constraints on the lower-level entity sets.

Representation of Generalization: Method 1

Method 1

- Create a schema for the higher-level entity set.
- For each lower-level entity set create it and link it with its upper-level schema.
- The primary-key attributes of the higher-level entity set become primary-key attributes of the higher-level entity set as well as all lower-level entity sets.
- We create foreign-key constraints on the lower-level entity sets.

Representation of Generalization: Method 1

Method 1

- Create a schema for the higher-level entity set.
- For each lower-level entity set create it and link it with its upper-level schema.
- The primary-key attributes of the higher-level entity set become primary-key attributes of the higher-level entity set as well as all lower-level entity sets.
- We create foreign-key constraints on the lower-level entity sets.

Representation of Generalization: Method 1

Method 1

- Create a schema for the higher-level entity set.
- For each lower-level entity set create it and link it with its upper-level schema.
- The primary-key attributes of the higher-level entity set become primary-key attributes of the higher-level entity set as well as all lower-level entity sets.
- We create foreign-key constraints on the lower-level entity sets.

Representation of Generalization: Method 1

Method 1

- Create a schema for the higher-level entity set.
- For each lower-level entity set create it and link it with its upper-level schema.
- The primary-key attributes of the higher-level entity set become primary-key attributes of the higher-level entity set as well as all lower-level entity sets.
- We create foreign-key constraints on the lower-level entity sets.

Representation of Generalization: Method 1 (Cont.)

Entities should look like:

person (ID , name, street, city)

employee (ID , salary)

student (ID , totalcredit)

Strength and Weakness of Method 1

Strength: Natural in design and reduces redundancies.

Weakness: Getting information requires joining 2 entities.



Representation of Generalization: Method 1 (Cont.)

Entities should look like:

person (ID , name, street, city)

employee (ID , salary)

student (ID , totalcredit)

Strength and Weakness of Method 1

Strength: Natural in design and reduces redundancies.

Weakness: Getting information requires joining 2 entities.



Representation of Generalization: Method 1 (Cont.)

Entities should look like:

person (ID , name, street, city)

employee (ID , salary)

student (ID , totalcredit)

Strength and Weakness of Method 1

Strength: Natural in design and reduces redundancies.

Weakness: Getting information requires joining 2 entities.

Representation of Generalization: Method 1 (Cont.)

Entities should look like:

person (ID , name, street, city)

employee (ID , salary)

student (ID , totalcredit)

Strength and Weakness of Method 1

Strength: Natural in design and reduces redundancies.

Weakness: Getting information requires joining 2 entities.



Representation of Generalization: Method 1 (Cont.)

Entities should look like:

person (ID , name, street, city)

employee (ID , salary)

student (ID , totalcredit)

Strength and Weakness of Method 1

Strength: Natural in design and reduces redundancies.

Weakness: Getting information requires joining 2 entities.



Representation of Generalization: Method 1 (Cont.)

Entities should look like:

person (ID , name, street, city)

employee (ID , salary)

student (ID , totalcredit)

Strength and Weakness of Method 1

Strength: Natural in design and reduces redundancies.

Weakness: Getting information requires joining 2 entities.



Representation of Generalization: Method 2

It can be applied if:

- If the generalization is **disjoint**
- If the generalization is **complete**

It implies:

If no entity is a member of two lower-level entity sets directly below a higher-level entity set, and if every entity in the higher-level entity set is also a member of one of the lower-level entity sets

Its ERD should look like:

employee (ID , name, street, city, salary)

student (ID , name, street, city, totalcredit)

Representation of Generalization: Method 2

It can be applied if:

- If the generalization is **disjoint**
- If the generalization is **complete**

It implies:

If no entity is a member of two lower-level entity sets directly below a higher-level entity set, and if every entity in the higher-level entity set is also a member of one of the lower-level entity sets

Its ERD should look like:

employee (ID , name, street, city, salary)

student (ID , name, street, city, totalcredit)

Representation of Generalization: Method 2

It can be applied if:

- If the generalization is **disjoint**
- If the generalization is **complete**

It implies:

If no entity is a member of two lower-level entity sets directly below a higher-level entity set, and if every entity in the higher-level entity set is also a member of one of the lower-level entity sets

Its ERD should look like:

employee (ID , name, street, city, salary)

student (ID , name, street, city, totalcredit)

Representation of Generalization: Method 2

It can be applied if:

- If the generalization is **disjoint**
- If the generalization is **complete**

It implies:

If no entity is a member of two lower-level entity sets directly below a higher-level entity set, and if every entity in the higher-level entity set is also a member of one of the lower-level entity sets

Its ERD should look like:

employee (ID , name, street, city, salary)

student (ID , name, street, city, totalcredit)

Representation of Generalization: Method 2

It can be applied if:

- If the generalization is **disjoint**
- If the generalization is **complete**

It implies:

If no entity is a member of two lower-level entity sets directly below a higher-level entity set, and if every entity in the higher-level entity set is also a member of one of the lower-level entity sets

Its ERD should look like:

employee (ID , name, street, city, salary)

student (ID , name, street, city, totalcredit)

Representation of Generalization: Method 2

It can be applied if:

- If the generalization is **disjoint**
- If the generalization is **complete**

It implies:

If no entity is a member of two lower-level entity sets directly below a higher-level entity set, and if every entity in the higher-level entity set is also a member of one of the lower-level entity sets

Its ERD should look like:

employee (ID , name, street, city, salary)

student (ID , name, street, city, totalcredit)

Method 2: Strength and weakness

- **Strength:** Data can be obtained directly, quick access (since no joining is done).
- **Weakness:** Data validation is failed (since no Foreign Key constraint can be imposed).
- **Weakness:** With **overlapping generalization**, some values would be stored multiple times, introduces **redundancies**. For instance, if a person is both an employee and a student, values for street and city would be stored twice.

Method 2: Strength and weakness

- **Strength:** Data can be obtained directly, quick access (since no joining is done).
- **Weakness:** Data validation is failed (since no Foreign Key constraint can be imposed).
- **Weakness:** With **overlapping generalization**, some values would be stored multiple times, introduces **redundancies**. For instance, if a person is both an employee and a student, values for street and city would be stored twice.

Method 2: Strength and weakness

- **Strength:** Data can be obtained directly, quick access (since no joining is done).
- **Weakness:** Data validation is failed (since no Foreign Key constraint can be imposed).
- **Weakness:** With **overlapping generalization**, some values would be stored multiple times, introduces **redundancies**. For instance, if a person is both an employee and a student, values for street and city would be stored twice.